

Московский Авиационный Институт

(Национальный Исследовательский Университет)

Институт №8 “Компьютерные науки и прикладная математика”

Кафедра №806 “Вычислительная математика и программирование”

Лабораторная работа №4 по курсу

«Операционные системы»

Группа: М8О-210Б-23

Студент: Воронухин Н.А.

Преподаватель: Бахарев В.Д.

Оценка: _____

Дата: 10.12.24

Москва, 2024

Постановка задачи

Вариант 1.

Требуется создать 2 динамические библиотеки реализующие два аллокатора памяти: списки свободных блоков (первый подходящий) и алгоритм двойников.

Общий метод и алгоритм решения

Использованные системные вызовы:

- `void exit(int status);` - приводит к обычному завершению программы, и величина `status` возвращается процессу-родителю.
- `void * mmap(void *start, size_t length, int prot, int flags, int fd, off_t offset);` - отражает `length` байтов, начиная со смещения `offset` файла (или другого объекта), определенного файловым дескриптором `fd`, в память, начиная с адреса `start`. Последний параметр (адрес) необязателен, и обычно бывает равен 0. Настоящее местоположение отраженных данных возвращается самой функцией `mmap`, и никогда не бывает равным 0.
- `int munmap(void *start, size_t length);` - удаляет отображение для указанного адресного диапазона
- `ssize_t write(int fd, const void buf[], size_t count);` - пишет `count` байт из буфера `buf` в файл, на который ссылается файловый дескриптор `fd`.
- `int mprotect(void *addr, size_t len, int prot);` - удаляет все отражения из заданной области памяти.
- `void *dlopen(const char *filename, int flag);` - загружает динамическую библиотеку, имя которой указано в строке `filename`, и возвращает прямой указатель на начало динамической библиотеки.
- `void *dlsym(void *handle, char *symbol);` - дает возможность получить по имени адрес символа, определенного в разделяемой библиотеке.
- `int dlclose(void *handle);` - закрывает динамическую библиотеку.

Реализованы 2 динамические библиотеки с аллокаторами памяти, первая со списком свободных блоков на основе односвязного списка, вторая – на основе алгоритма двойников с помощью массива списков, хранящих блоки разных размеров, и `main`-файл для их проверки.

Код программы

dl.h

```
#ifndef __LIBRARY_H
```

```

#define __LIBRARY_H

#include <dlfcn.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
#include <dlfcn.h>
#include <sys/mman.h>
#include <stddef.h>
#include <string.h>
#include <stdbool.h>

#ifdef _MSC_VER
#define EXPORT __declspec(dllexport)
#else
#define EXPORT
#endif

typedef struct Allocator Allocator;
typedef struct Block Block;

typedef Allocator *allocator_create_f(void *const memory, const size_t size);
typedef void allocator_destroy_f(Allocator *const allocator);
typedef void *allocator_alloc_f(Allocator *const allocator, const size_t size);
typedef void allocator_free_f(Allocator *const allocator, void *const memory);

#endif

```

parent.c

```

#include "dl.h"

```

```
#include <sys/mman.h>
```

```
#define MEMORY_POOL_SIZE 1024 * 1024
```

```
static Allocator *allocator_create_stub(void *const memory, const size_t size)
```

```
{
```

```
    void *mapped_memory = mmap(memory, size, PROT_READ | PROT_WRITE, MAP_PRIVATE |  
MAP_ANONYMOUS | MAP_FIXED, -1, 0);
```

```
    if (mapped_memory == MAP_FAILED)
```

```
    {
```

```
        const char err_msg[] = "allocator_create: mmap failed\n";
```

```
        write(STDERR_FILENO, err_msg, sizeof(err_msg) - 1);
```

```
        return NULL;
```

```
    }
```

```
    return (Allocator *)mapped_memory;
```

```
}
```

```
static void allocator_destroy_stub(Allocator *const allocator)
```

```
{
```

```
    if (allocator && munmap(allocator, MEMORY_POOL_SIZE) == -1)
```

```
    {
```

```
        const char err_msg[] = "allocator_destroy: munmap failed\n";
```

```
        write(STDERR_FILENO, err_msg, sizeof(err_msg) - 1);
```

```
    }
```

```
}
```

```
static void *allocator_alloc_stub(Allocator *const allocator, const size_t size)
```

```
{
```

```
    void *mapped_memory = mmap(NULL, size, PROT_READ | PROT_WRITE, MAP_PRIVATE |  
MAP_ANONYMOUS, -1, 0);
```

```
    if (mapped_memory == MAP_FAILED)
```

```
    {
```

```
        const char err_msg[] = "allocator_alloc: mmap failed\n";
```

```

        write(STDERR_FILENO, err_msg, sizeof(err_msg) - 1);

        return NULL;
    }

    return mapped_memory;
}

static void allocator_free_stub(Allocator *const allocator, void *const memory)
{
    if (memory && munmap(memory, sizeof(memory)) == -1)
    {
        const char err_msg[] = "allocator_free: munmap failed\n";
        write(STDERR_FILENO, err_msg, sizeof(err_msg) - 1);
    }
}

static allocator_create_f *allocator_create;
static allocator_destroy_f *allocator_destroy;
static allocator_alloc_f *allocator_alloc;
static allocator_free_f *allocator_free;

int main(int argc, char **argv)
{
    void *library = NULL;

    if (argc == 2) {
        library = dlopen(argv[1], RTLD_LOCAL | RTLD_NOW);

        if (!library)
        {
            const char msg[] = "Failed to load library. Using stub.\n";
            write(STDERR_FILENO, msg, sizeof(msg));

            allocator_create = allocator_create_stub;
            allocator_destroy = allocator_destroy_stub;

```

```

    allocator_alloc = allocator_alloc_stub;

    allocator_free = allocator_free_stub;
}
else
{
    allocator_create = dlsym(library, "allocator_create");
    allocator_destroy = dlsym(library, "allocator_destroy");
    allocator_alloc = dlsym(library, "allocator_alloc");
    allocator_free = dlsym(library, "allocator_free");

    if (!allocator_create)
    {
        const char msg[] = "Failed to load allocator_create. Using stub.\n";
        write(STDERR_FILENO, msg, sizeof(msg));

        allocator_create = allocator_create_stub;
    }
    if (!allocator_destroy)
    {
        const char msg[] = "Failed to load allocator_destroy. Using stub.\n";
        write(STDERR_FILENO, msg, sizeof(msg));

        allocator_destroy = allocator_destroy_stub;
    }
    if (!allocator_alloc)
    {
        const char msg[] = "Failed to load allocator_alloc. Using stub.\n";
        write(STDERR_FILENO, msg, sizeof(msg));

        allocator_alloc = allocator_alloc_stub;
    }
    if (!allocator_free)
    {
        const char msg[] = "Failed to load allocator_free. Using stub.\n";
        write(STDERR_FILENO, msg, sizeof(msg));

        allocator_free = allocator_free_stub;
    }
}

```

```

        }
    }
}
else
{
    const char msg[] = "Incorrect use. Using stub.\n";
    write(STDERR_FILENO, msg, sizeof(msg));
    allocator_create = allocator_create_stub;
    allocator_destroy = allocator_destroy_stub;
    allocator_alloc = allocator_alloc_stub;
    allocator_free = allocator_free_stub;
}

// TeCfC,C< P±PëP±P»PëPsC,PμPePë
size_t size = MEMORY_POOL_SIZE;
void *addr = mmap(NULL, size, PROT_READ | PROT_WRITE,
                  MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
if (addr == MAP_FAILED)
{
    dlclose(library);
    char message[] = "mmap failed\n";
    write(STDERR_FILENO, message, sizeof(message) - 1);
    return EXIT_FAILURE;
}

Allocator *allocator = allocator_create(addr, MEMORY_POOL_SIZE);

if (!allocator)
{
    const char msg[] = "Failed to initialize allocator\n";
    write(STDERR_FILENO, msg, sizeof(msg));
    munmap(addr, size);
    dlclose(library);
}

```

```

        return EXIT_FAILURE;
    }

    int *int_block = (int *)allocator_alloc(allocator, sizeof(int));
    if (int_block)
    {
        *int_block = 42;

        const char msg[] = "Allocated int_block with value 42\n";
        write(STDOUT_FILENO, msg, sizeof(msg));
    }
    else
    {
        const char msg[] = "Failed to allocate memory for int_block\n";
        write(STDERR_FILENO, msg, sizeof(msg));
    }

    float *float_block = (float *)allocator_alloc(allocator, sizeof(float));
    if (float_block)
    {
        *float_block = 2.718f;

        const char msg[] = "Allocated float_block with value 2.718\n";
        write(STDOUT_FILENO, msg, sizeof(msg));
    }
    else
    {
        const char msg[] = "Failed to allocate memory for float_block\n";
        write(STDERR_FILENO, msg, sizeof(msg));
    }

    if (int_block)
    {
        allocator_free(allocator, int_block);
    }

```



```

        const char msg[] = "Freed int_block\n";
        write(STDOUT_FILENO, msg, sizeof(msg));
    }

    if (float_block)
    {
        allocator_free(allocator, float_block);
        const char msg[] = "Freed float_block\n";
        write(STDOUT_FILENO, msg, sizeof(msg));
    }

    allocator_destroy(allocator);
    const char msg[] = "Allocator destroyed\n";
    write(STDOUT_FILENO, msg, sizeof(msg));

    if (library)
        dlclose(library);
    munmap(addr, size);

    return EXIT_SUCCESS;
}

```

childe.c

```

#include "dl.h"
#include <math.h>

#define MIN_BLOCK_SIZE 32
#define MAX_BLOCK_INDEX 24

typedef struct Block
{
    size_t size;
    struct Block *next;
}

```

```
} Block;
```

```
typedef struct Allocator
```

```
{
```

```
    void *memory;
```

```
    size_t total_size;
```

```
    Block *free_lists[MAX_BLOCK_INDEX];
```

```
} Allocator;
```

```
size_t round_to_power_of_two(size_t size)
```

```
{
```

```
    size_t power = MIN_BLOCK_SIZE;
```

```
    while (power < size)
```

```
    {
```

```
        power <<= 1;
```

```
    }
```

```
    return power;
```

```
}
```

```
int get_power_of_two(size_t size)
```

```
{
```

```
    return (int)(log2(size));
```

```
}
```

```
EXPORT Allocator *allocator_create(void *memory, const size_t size)
```

```
{
```

```
    if (!memory || size < MIN_BLOCK_SIZE + sizeof(Allocator) + sizeof(Block))
```

```
    {
```

```
        return NULL;
```

```
    }
```

```
    Allocator *allocator = (Allocator*)memory; //(Allocator *)mmap(
```

```
        //NULL, sizeof(Allocator), PROT_READ | PROT_WRITE, MAP_PRIVATE |  
MAP_ANONYMOUS, -1, 0);
```

```
    size_t total_size = round_to_power_of_two(size - sizeof(Allocator));
```

```
    size_t max_size = 1 << (MAX_BLOCK_INDEX - 1);
```

```
    if (total_size > max_size)
```

```
    {
```

```
        return NULL;
```

```
    }
```

```
    allocator->memory = memory - sizeof(Allocator);
```

```
    allocator->total_size = total_size;
```

```
    for (size_t i = 0; i < MAX_BLOCK_INDEX; i++)
```

```
    {
```

```
        allocator->free_lists[i] = NULL;
```

```
    }
```

```
    Block *initial_block = (Block *)allocator->memory;
```

```
    initial_block->size = total_size;
```

```
    initial_block->next = NULL;
```

```
    int index = get_power_of_two(total_size) - get_power_of_two(MIN_BLOCK_SIZE);
```

```
    allocator->free_lists[index] = initial_block;
```

```
    return allocator;
```

```
}
```

```
EXPORT void allocator_destroy(Allocator *allocator)
```

```
{
```

```
    if (!allocator)
```

```
        return;
```

```

    memset(allocator, 0, allocator->total_size + sizeof(Allocator));
}

EXPORT void *allocator_alloc(Allocator *const allocator, const size_t size)
{
    if (!allocator || size == 0)
        return NULL;

    size_t block_size = round_to_power_of_two(size + sizeof(Block));
    size_t max_size = 1 << (MAX_BLOCK_INDEX - 1);

    if (block_size > max_size)
    {
        return NULL;
    }

    int index = get_power_of_two(block_size) - get_power_of_two(MIN_BLOCK_SIZE);
    if (index >= MAX_BLOCK_INDEX)
    {
        return NULL;
    }
    else if (index < 0)
    {
        index = 1;
    }

    while (index < MAX_BLOCK_INDEX && !allocator->free_lists[index])
    {
        index++;
    }

    if (index >= MAX_BLOCK_INDEX)
    {

```

```
    return NULL;
}
```

```
Block *block = allocator->free_lists[index];
allocator->free_lists[index] = block->next;
```

```
while (block->size > block_size)
```

```
{
    size_t new_size = block->size >> 1;
    Block *buddy = (Block *)((char *)block + new_size);
    buddy->size = new_size;
    buddy->next =
```

```
        allocator->free_lists[get_power_of_two(new_size) -
get_power_of_two(MIN_BLOCK_SIZE)];
```

```
        allocator->free_lists[get_power_of_two(new_size) -
get_power_of_two(MIN_BLOCK_SIZE)] =
```

```
            buddy;
    block->size = new_size;
}
```

```
return (void *)((char *)block + sizeof(Block));
}
```

```
EXPORT void allocator_free(Allocator *const allocator, void *const memory)
```

```
{
    if (!allocator || !memory)
        return;
```

```
Block *block = (Block *)((char *)memory - sizeof(Block));
```

```
size_t block_size = block->size;
```

```
int index = get_power_of_two(block_size) - get_power_of_two(MIN_BLOCK_SIZE);
```

```
if (index >= MAX_BLOCK_INDEX)
    return;
```

```
        block->next = allocator->free_lists[index];

        allocator->free_lists[index] = block;
    }
}
```

free_list.c

```
#include "dl.h"
```

```
#define MIN_BLOCK_SIZE 32
```

```
typedef struct Block
```

```
{
    size_t size;
    struct Block *next;
    bool is_free;
} Block;
```

```
typedef struct Allocator
```

```
{
    Block *free_list;
    void *memory_start;
    size_t total_size;
} Allocator;
```

```
EXPORT Allocator *allocator_create(void *memory, size_t size)
```

```
{
    if (!memory || size < sizeof(Allocator) + sizeof(Block) + MIN_BLOCK_SIZE)
    {
        return NULL;
    }
}
```

```
Allocator *allocator = (Allocator *)memory;
allocator->memory_start = (char *)memory + sizeof(Allocator);
allocator->total_size = size - sizeof(Allocator);
```

```

allocator->free_list = (Block *)allocator->memory_start;

allocator->free_list->size = allocator->total_size - sizeof(Block);
allocator->free_list->next = NULL;
allocator->free_list->is_free = true;

return allocator;
}

EXPORT void allocator_destroy(Allocator *allocator)
{
    if (allocator)
    {
        memset(allocator, 0, allocator->total_size + sizeof(Allocator));
    }
}

EXPORT void *allocator_alloc(Allocator *allocator, size_t size)
{
    if (!allocator || size < 1)
    {
        return NULL;
    }

    size = (size / MIN_BLOCK_SIZE
            + (size % MIN_BLOCK_SIZE) ? 1 : 0)
        * MIN_BLOCK_SIZE;

    Block *current = allocator->free_list;

    while (current)
    {
        if (current->is_free && current->size >= size) {

```

```

    size_t remain_size = current->size - size;

    if (remain_size >= sizeof(Block) + MIN_BLOCK_SIZE) {

        Block *new_block = (Block *)((char *)current + 2 * sizeof(Block) + size);

        new_block->size = remain_size - sizeof(Block);

        new_block->is_free = true;

        new_block->next = current->next;

        current->next = new_block;

        current->size = size;

    }

    current->is_free = false;

    return (void *)((char *)current + sizeof(Block));

}

else {

    current = current->next;

}

}

return NULL;

}

```

```

EXPORT void allocator_free(Allocator *allocator, void *ptr_to_memory)

{

    if (!allocator || !ptr_to_memory) // bound

    {

        return;

    }

    Block *current = (Block *)((char *)ptr_to_memory - sizeof(Block));

    if (!current)

        return;

    current->is_free = true;

```



```
allocator->free_list;
while (current && current->next)
{
    if (current->is_free && current->next->is_free)
    {
        current->size += current->next->size + sizeof(Block);
        current->next = current->next->next;
    }
    else
    {
        current = current->next;
    }
}
}
```

Результаты тестирования

Метод свободных блоков.

- **Производительность:** Быстрое выделение памяти, но медленное объединение блоков.
- **Фрагментация:** присутствует как внутренняя, так и внешняя фрагментации, внутренняя зависит от размеров минимального блока, внешняя от последовательности запросов памяти.
- **Память:** Эффективен для запросов практически любого размера

Метод двойников.

- **Производительность:** Высокая скорость выделение памяти, и освобождения.
- **Фрагментация:** присутствует как сильная внутренняя, так и внешняя фрагментации, внутренняя зависит от размеров минимального блока и кратности степени 2 размеров запросов, внешняя от последовательности запросов памяти.
- **Память:** Эффективен для запросов кратных степени 2.

Протокол работы программы

```
vnadez@dbnicknv:~/Projects/OS_LABS/LAB_4$ ./main  
/home/vnadez/Projects/OS_LABS/LAB_4/bd.so
```

```
Allocated int_block with value 42
```

```
Allocated float_block with value 2.718
```

```
Freed int_block
```

```
Freed float_block
```

```
Allocator destroyed
```

```
vnadez@dbnicknv:~/Projects/OS_LABS/LAB_4$ ./main
```

```
Incorrect use. Using stub.
```

```
Allocated int_block with value 42
```

```
Allocated float_block with value 2.718
```

```
Freed int_block
```

```
Freed float_block
```

```
Allocator destroyed
```

Strace:

```
16463 execve("./main", [ "./main", "/home/vnadez/Projects/OS_LABS/LA"... ],  
0x7fff74df1b90 /* 64 vars */) = 0
```

```

16463 brk(NULL) = 0x560e4d62c000

16463 mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7ff8977cc000

16463 access("/etc/ld.so.preload", R_OK) = -1 ENOENT (Нет такого файла или каталога)

16463 openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3

16463 newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=97022, ...}, AT_EMPTY_PATH) = 0

16463 mmap(NULL, 97022, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7ff8977b4000

16463 close(3) = 0

16463 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3

16463 read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\20t\2\0\0\0\0"...
832) = 832

16463 pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0"...
784, 64) = 784

16463 newfstatat(3, "", {st_mode=S_IFREG|0755, st_size=1922136, ...}, AT_EMPTY_PATH) =
0

16463 pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0"...
784, 64) = 784

16463 mmap(NULL, 1970000, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7ff8975d3000

16463 mmap(0x7ff8975f9000, 1396736, PROT_READ|PROT_EXEC,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x26000) = 0x7ff8975f9000

16463 mmap(0x7ff89774e000, 339968, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3,
0x17b000) = 0x7ff89774e000

16463 mmap(0x7ff8977a1000, 24576, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1ce000) = 0x7ff8977a1000

16463 mmap(0x7ff8977a7000, 53072, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7ff8977a7000

16463 close(3) = 0

16463 mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7ff8975d0000

16463 arch_prctl(ARCH_SET_FS, 0x7ff8975d0740) = 0

16463 set_tid_address(0x7ff8975d0a10) = 16463

16463 set_robust_list(0x7ff8975d0a20, 24) = 0

16463 rseq(0x7ff8975d1060, 0x20, 0, 0x53053053) = 0

16463 mprotect(0x7ff8977a1000, 16384, PROT_READ) = 0

16463 mprotect(0x560e44fc1000, 4096, PROT_READ) = 0

16463 mprotect(0x7ff8977fe000, 8192, PROT_READ) = 0

16463 prlimit64(0, RLIMIT_STACK, NULL, {rlim_cur=8192*1024, rlim_max=RLIM64_INFINITY})
= 0

16463 munmap(0x7ff8977b4000, 97022) = 0

16463 getrandom("\x51\x50\xa6\x2b\x81\x1f\x3d\xcf", 8, GRND_NONBLOCK) = 8

```

```

16463 brk(NULL) = 0x560e4d62c000

16463 brk(0x560e4d64d000) = 0x560e4d64d000

16463 openat(AT_FDCWD, "/home/vnadez/Projects/OS_LABS/LAB_4/fr.so",
O_RDONLY|O_CLOEXEC) = 3

16463 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0\0\0\0\0"...
832) = 832

16463 newfstatat(3, "", {st_mode=S_IFREG|0755, st_size=15440, ...}, AT_EMPTY_PATH) = 0

16463 mmap(NULL, 16408, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7ff8977c7000

16463 mmap(0x7ff8977c8000, 4096, PROT_READ|PROT_EXEC,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1000) = 0x7ff8977c8000

16463 mmap(0x7ff8977c9000, 4096, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3,
0x2000) = 0x7ff8977c9000

16463 mmap(0x7ff8977ca000, 8192, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x2000) = 0x7ff8977ca000

16463 close(3) = 0

16463 mprotect(0x7ff8977ca000, 4096, PROT_READ) = 0

16463 mmap(NULL, 1048576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7ff8974d0000

16463 write(1, "Allocated int_block with value 4"... , 35) = 35

16463 write(1, "Allocated float_block with value"... , 40) = 40

16463 write(1, "Freed int_block\n\0", 17) = 17

16463 write(1, "Freed float_block\n\0", 19) = 19

16463 write(1, "Allocator destroyed\n\0", 21) = 21

16463 munmap(0x7ff8977c7000, 16408) = 0

16463 munmap(0x7ff8974d0000, 1048576) = 0

16463 exit_group(0) = ?

16463 +++ exited with 0 +++

```

Вывод

В процессе выполнения лабораторной работы я освоил работу с динамическими библиотеками, новыми системными вызовами, предназначенными для работы с ними, написание собственного аллокатора памяти на языке C, а также проанализировал работу пары алгоритмов выделения памяти, самым проблемным из которых для меня оказался метод двойников.